Ajuntament de Barcelona

Data Massage - R

Introduction Data Massage with R

IT Academy

Barcelona Activa

Unió Europea
Fons Europeu de Desenvolupament Regional

Index

- Data massaging with R
- What is tidyr?
- What is dplyr?

Data massaging with R

Data Massaging is the process of extracting data to remove unneeded information or clean up a dataset to get into a usable format.

- To do it in R, we can use **dplyr** and **tidyr** packages

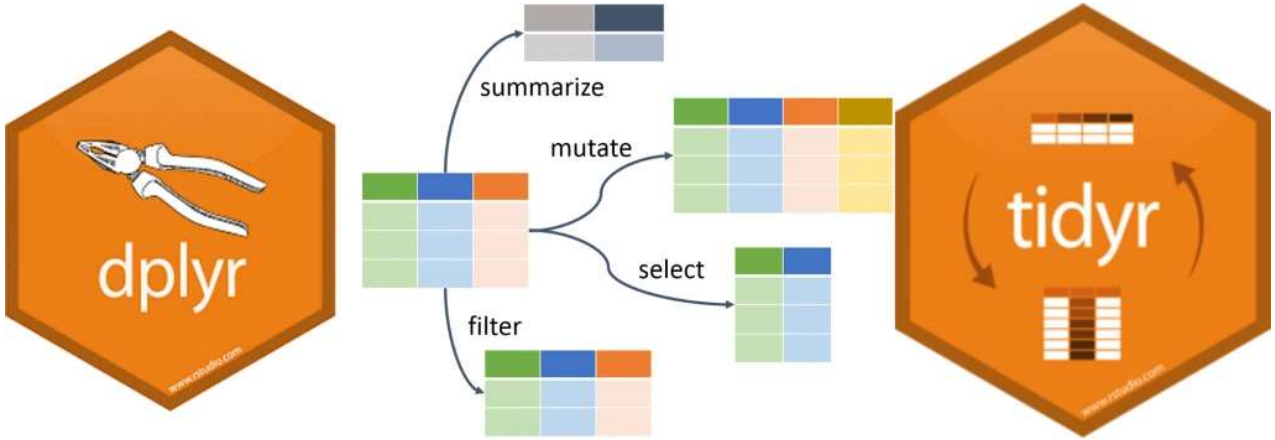

What is tidyr?

tidyr is a package that makes it easy to "tidy" your data.

Tidy data is data that's easy to work with: it's easy to munge (with **dplyr**), visualise (with **ggplot2** or ggvis) and model (with R's hundreds of modelling packages). The two most important properties of tidy data are:

- Each variable forms a column
- Each observation forms a row

Tidy data graphic from R for Data Science

What is dplyr?

dplyr is a package which provides a set of tools for efficiently manipulating datasets in R

You can track developement progress at https://github.com/tidyverse/dplyr

dplyr

Overview

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

These all combine naturally with `group_by()` which allows you to perform any operation "by group". You can learn more about them in `vignette("dplyr")`. As well as these single-table verbs, dplyr also provides a variety of two-table verbs, which you can learn about in `vignette("two-table")`.

Basics from tidyr and dplyr

Package	Function	Use
dplyr	select	select variables/columns
dplyr	filter	select observations/rows
dplyr	Mutate	transform or recode variables
dplyr	summarize / summarise	summarize data
dplyr	group_by	identify subgroups for further processing
tidyr	gather	convert wide format dataset to long format
tidyr	spread	convert long format dataset to wide format

You are going to use:

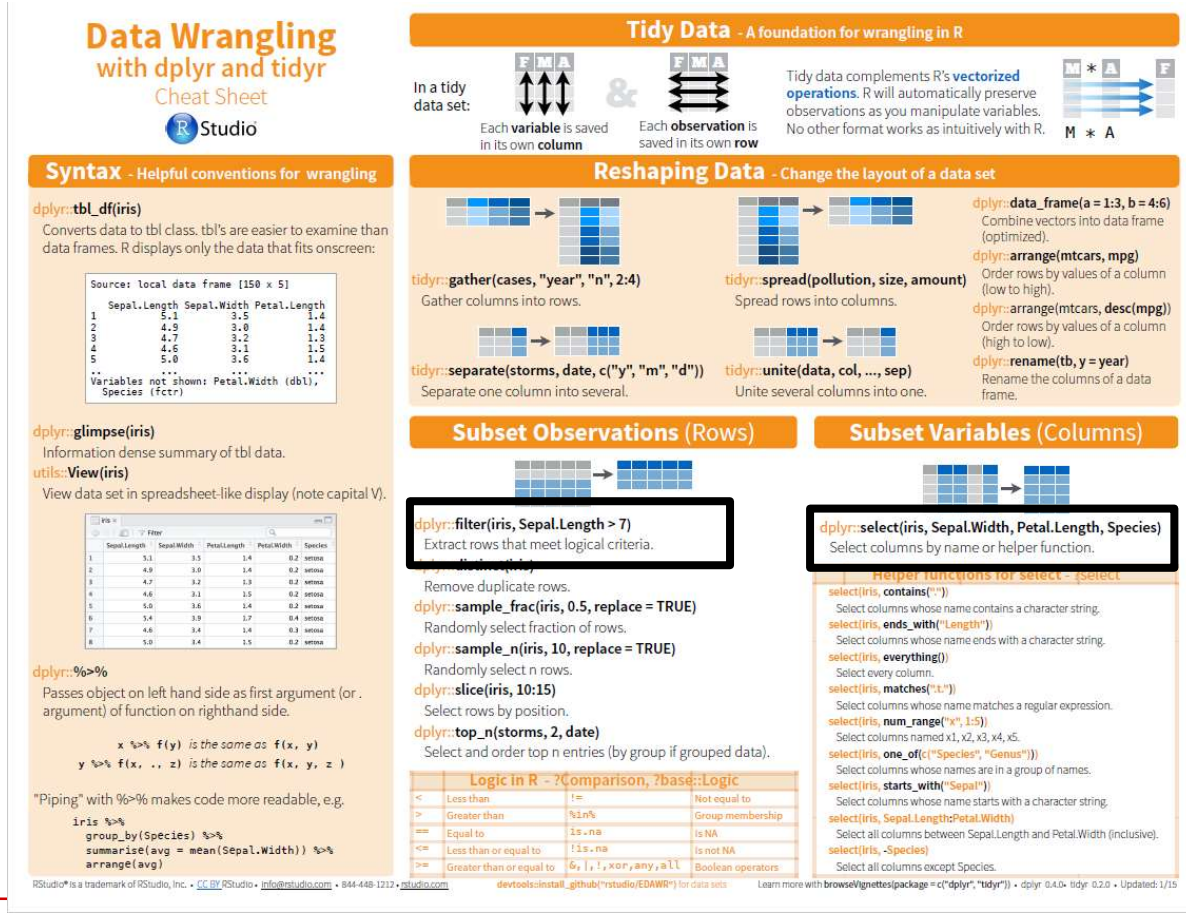

Data Wrangling with dplyr and tidyr Cheat Sheet

RStudio

Syntax - Helpful conventions for wrangling

dplyr::tbl_df(iris)
Converts data to tbl class. tbl's are easier to examine than data frames. R displays only the data that fits onscreen:

```
Source: local data frame [150 x 5]

   Sepal.Length Sepal.Width Petal.Length
1           5.1         3.5          1.4
2           4.9         3.0          1.4
3           4.7         3.2          1.3
4           4.6         3.1          1.5
5           5.0         3.6          1.4
..          ...         ...          ...
Variables not shown: Petal.Width (dbl),
  Species (fctr)
```

dplyr::glimpse(iris)
Information dense summary of tbl data.

utils::View(iris)
View data set in spreadsheet-like display (note capital V).

dplyr::%>%
Passes object on left hand side as first argument (or . argument) of function on righthand side.

```
x %>% f(y)  is the same as  f(x, y)
y %>% f(x, .., z)  is the same as  f(x, y, z)
```

"Piping" with %>% makes code more readable, e.g.

```
iris %>%
  group_by(Species) %>%
  summarise(avg = mean(Sepal.Width)) %>%
  arrange(avg)
```

Tidy Data - A foundation for wrangling in R

In a tidy data set:

Each **variable** is saved in its own **column**

Each **observation** is saved in its own **row**

Tidy data complements R's **vectorized operations**. R will automatically preserve observations as you manipulate variables. No other format works as intuitively with R.

M * A → F

Reshaping Data - Change the layout of a data set

tidyr::gather(cases, "year", "n", 2:4)
Gather columns into rows.

tidyr::spread(pollution, size, amount)
Spread rows into columns.

tidyr::separate(storms, date, c("y", "m", "d"))
Separate one column into several.

tidyr::unite(data, col, ..., sep)
Unite several columns into one.

dplyr::data_frame(a = 1:3, b = 4:6)
Combine vectors into data frame (optimized).

dplyr::arrange(mtcars, mpg)
Order rows by values of a column (low to high).

dplyr::arrange(mtcars, desc(mpg))
Order rows by values of a column (high to low).

dplyr::rename(tb, y = year)
Rename the columns of a data frame.

Subset Observations (Rows)

dplyr::filter(iris, Sepal.Length > 7)
Extract rows that meet logical criteria.

dplyr::distinct(iris)
Remove duplicate rows.

dplyr::sample_frac(iris, 0.5, replace = TRUE)
Randomly select fraction of rows.

dplyr::sample_n(iris, 10, replace = TRUE)
Randomly select n rows.

dplyr::slice(iris, 10:15)
Select rows by position.

dplyr::top_n(storms, 2, date)
Select and order top n entries (by group if grouped data).

Logic in R - ?Comparison, ?base::Logic

<	Less than	!=	Not equal to
>	Greater than	%in%	Group membership
==	Equal to	is.na	is NA
<=	Less than or equal to	!is.na	is not NA
>=	Greater than or equal to	&, |, !, xor, any, all	Boolean operators

Subset Variables (Columns)

dplyr::select(iris, Sepal.Width, Petal.Length, Species)
Select columns by name or helper function.

Helper functions for select - ?select

select(iris, contains("."))
Select columns whose name contains a character string.

select(iris, ends_with("Length"))
Select columns whose name ends with a character string.

select(iris, everything())
Select every column.

select(iris, matches("t."))
Select columns whose name matches a regular expression.

select(iris, num_range("x", 1:5))
Select columns named x1, x2, x3, x4, x5.

select(iris, one_of(c("Species", "Genus")))
Select columns whose names are in a group of names.

select(iris, starts_with("Sepal"))
Select columns whose name starts with a character string.

select(iris, Sepal.Length:Petal.Width)
Select all columns between Sepal.Length and Petal.Width (inclusive).

select(iris, -Species)
Select all columns except Species.

devtools::install_github("rstudio/EDAWR") for data sets

Learn more with browseVignettes(package = c("dplyr", "tidyr")) • dplyr 0.4.0• tidyr 0.2.0 • Updated: 1/15

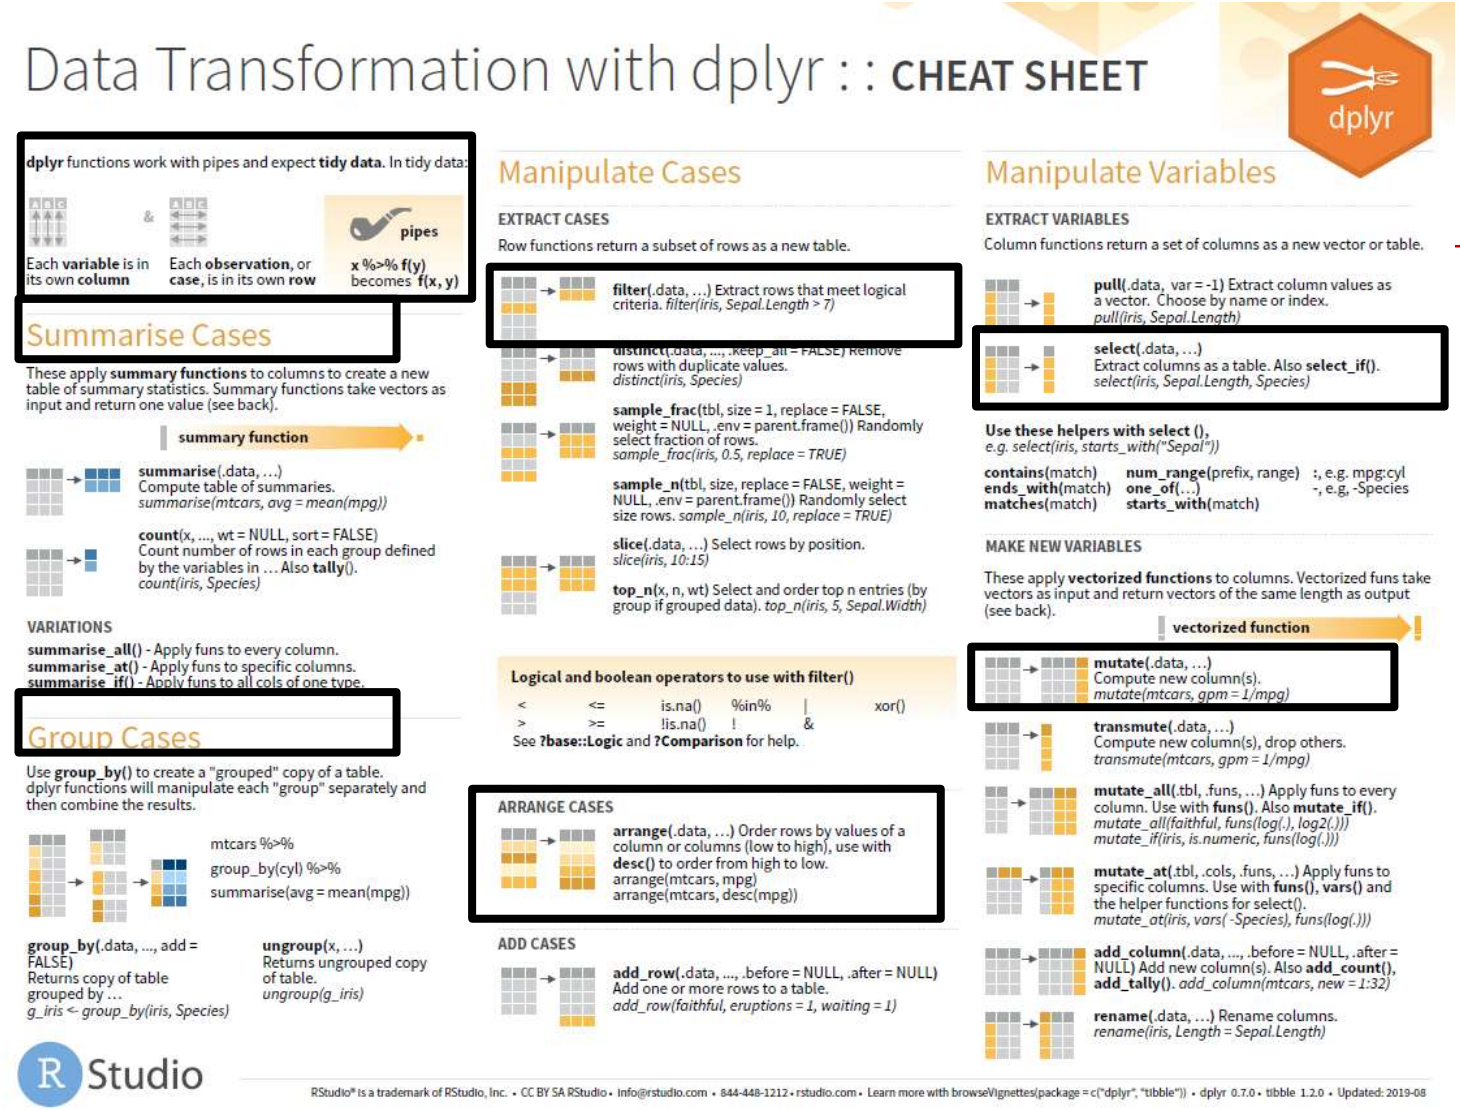
Data Transformation with dplyr : : CHEAT SHEET
dplyr
dplyr functions work with pipes and expect tidy data. In tidy data:
Each variable is in its own column
Each observation, or case, is in its own row
pipes
x %>% f(y) becomes f(x, y)
Summarise Cases
These apply summary functions to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).
summary function
summarise(.data, …) Compute table of summaries.
summarise(mtcars, avg = mean(mpg))
count(x, ..., wt = NULL, sort = FALSE) Count number of rows in each group defined by the variables in … Also tally().
count(iris, Species)
VARIATIONS
summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.
Group Cases
Use group_by() to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.
mtcars %>%
group_by(cyl) %>%
summarise(avg = mean(mpg))
group_by(.data, ..., add = FALSE) Returns copy of table grouped by …
g_iris <- group_by(iris, Species)
ungroup(x, …) Returns ungrouped copy of table.
ungroup(g_iris)
Manipulate Cases
EXTRACT CASES
Row functions return a subset of rows as a new table.
filter(.data, …) Extract rows that meet logical criteria. filter(iris, Sepal.Length > 7)
distinct(.data, ..., .keep_all = FALSE) Remove rows with duplicate values.
distinct(iris, Species)
sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, .env = parent.frame()) Randomly select fraction of rows.
sample_frac(iris, 0.5, replace = TRUE)
sample_n(tbl, size, replace = FALSE, weight = NULL, .env = parent.frame()) Randomly select size rows. sample_n(iris, 10, replace = TRUE)
slice(.data, …) Select rows by position.
slice(iris, 10:15)
top_n(x, n, wt) Select and order top n entries (by group if grouped data). top_n(iris, 5, Sepal.Width)
Logical and boolean operators to use with filter()
< <= is.na() %in% | xor()
> >= !is.na() ! &
See ?base::Logic and ?Comparison for help.
ARRANGE CASES
arrange(.data, …) Order rows by values of a column or columns (low to high), use with desc() to order from high to low.
arrange(mtcars, mpg)
arrange(mtcars, desc(mpg))
ADD CASES
add_row(.data, ..., .before = NULL, .after = NULL) Add one or more rows to a table.
add_row(faithful, eruptions = 1, waiting = 1)
Manipulate Variables
EXTRACT VARIABLES
Column functions return a set of columns as a new vector or table.
pull(.data, var = -1) Extract column values as a vector. Choose by name or index.
pull(iris, Sepal.Length)
select(.data, …) Extract columns as a table. Also select_if().
select(iris, Sepal.Length, Species)
Use these helpers with select (),
e.g. select(iris, starts_with("Sepal"))
contains(match) num_range(prefix, range) :, e.g. mpg:cyl
ends_with(match) one_of(…) -, e.g, -Species
matches(match) starts_with(match)
MAKE NEW VARIABLES
These apply vectorized functions to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).
vectorized function
mutate(.data, …) Compute new column(s).
mutate(mtcars, gpm = 1/mpg)
transmute(.data, …) Compute new column(s), drop others.
transmute(mtcars, gpm = 1/mpg)
mutate_all(.tbl, .funs, …) Apply funs to every column. Use with funs(). Also mutate_if().
mutate_all(faithful, funs(log(.), log2(.)))
mutate_if(iris, is.numeric, funs(log(.)))
mutate_at(.tbl, .cols, .funs, …) Apply funs to specific columns. Use with funs(), vars() and the helper functions for select().
mutate_at(iris, vars(-Species), funs(log(.)))
add_column(.data, ..., .before = NULL, .after = NULL) Add new column(s). Also add_count(), add_tally(). add_column(mtcars, new = 1:32)
rename(.data, …) Rename columns.
rename(iris, Length = Sepal.Length)
RStudio
RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • Info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with browseVignettes(package = c("dplyr", "tibble")) • dplyr 0.7.0 • tibble 1.2.0 • Updated: 2019-08

Barcelona Activa

30 ANYS

barcelona.cat/barcelonactiva